

You are a data scientist hired by TelCoX, a telecommunications company struggling with high customer churn rates. Your mission is to build a logistic regression model to predict customer churn and provide actionable insights to reduce churn. You are given access to a Kaggle dataset containing historical customer data.

1. Data Preprocessing:

- a. Load the TelCoX Customer Churn Dataset and inspect its structure.
- b. Handle missing values appropriately.
- c. Encode categorical variables for modeling.

2. Logistic Regression Model Building:

- a. Split the dataset into training and testing sets (80% training, 20% testing).
- b. Build a baseline logistic regression model using default settings.
- c. Train the model on the training data.

3. Model Evaluation and Visualization:

- a. Evaluate the model's performance on the testing data using the following metrics:
 - Accuracy
 - Precision
 - Recall
 - F1-score
 - ROC AUC
 - Cohen's Kappa Score
 - Matthews Correlation Coefficient
- b. Create a confusion matrix visualization.
- c. Plot the ROC curve and calculate the AUC.
- d. Visualize the precision-recall curve.
- e. Plot the log loss function curve during model training.

4. Sigmoid Curve:

- a. Provide a clear visualization of the sigmoid (logistic) curve used in logistic regression. Explain its significance.

5. Hyperparameter Tuning:

- a. Perform hyperparameter tuning to optimize the logistic regression model. Experiment with different values for at least one hyperparameter.
- b. Discuss how hyperparameter tuning affects the model's performance.

6. Recommendations:

- a. Based on your analysis, provide actionable recommendations to TelCoX on reducing customer churn.